

Task: Build a Voice AI Agent

Objective

Design and implement a voice AI agent capable of making outbound calls for a specific scenario. The primary evaluation focus is your solution architecture, how you connect speech, language, and telephony components into a working, production-minded system.

Requirements

Minimal UI

Build a simple interface where a user can:

- Enter a phone number and select or define a call scenario
- Trigger an outbound call to that number

UI can be as simple as you want, a clean single-page layout is perfectly fine.

Backend

All backend logic must be built using FastAPI.

Note:

If you are shortlisted for the next round, there may be additional changes or revisions required in the current task. Make sure to keep the implementation as dynamic and flexible as possible.

Functionality

Choose one realistic outbound campaign scenario to implement. Examples:

- Appointment reminder and confirmation
- Lead qualification call
- Customer satisfaction survey
- Payment follow-up
- Event registration confirmation

Design a clear agent persona, script, and conversation flow for your chosen scenario. The agent should handle basic user responses naturally.

Voice AI Approach:

Use a platform that handles STT, LLM, and TTS end-to-end like Vapi, Retell AI, Bland AI or

Build the pipeline yourself using individual services:

- STT: Deepgram, Whisper, AssemblyAI, or similar
- LLM: OpenAI, Claude, Groq, Together AI, or open-source
- TTS: ElevenLabs, OpenAI TTS, Google TTS, or similar
- Telephony: Twilio, Vonage, or platform-provided

Deliverables

- GitHub repository with full source code (Working)
- README with setup instructions, architecture overview, and your design decisions

Evaluation Focus

- Clean implementation
- FastAPI structure & design
- Solution architecture
- Context aware conversation